

YUZHE MA

University of Wisconsin–Madison
ma234@wisc.edu · 6082133287

Education

University of Wisconsin–Madison Ph.D. in Computer Sciences, Minor in Statistics Advisor: Professor Xiaojin (Jerry) Zhu	05/2021 (<i>expected</i>)
University of Wisconsin–Madison M.S. in Computer Sciences	05/2018
Huazhong University of Science and Technology B.S. in Computer Science and Technology	06/2016

Research Interests

My research interest lies in machine learning, especially adversarial machine learning and sequential decision making such as multi-armed bandit and reinforcement learning. I am also interested in solving real-world problems related to recommendation systems, search ranking, deep learning, unsupervised learning (e.g., dimensionality reduction and clustering), differential privacy, trust-worthy AI, machine teaching, and optimization.

Industrial Experience

Research Intern, IBM Research ◦ Built a two-stage machine learning model based on Cycle Generative Adversarial Network (Cycle-GAN) and object detection techniques to identify buildings on the satellite imagery data of American cities. ◦ Developed an iterative training procedure based on the object detection algorithm YOLO to augment the labels in the training data. The model increased the total amount of labeled buildings from 20K to 80K. ◦ Applied Cycle-GAN to transform imagery data into rasterized maps.	06-09/2020
Applied Scientist Intern, Amazon ◦ Developed a student identification model using the Gradient Boosted Tree (GBT) algorithm. ◦ Carried out an end-to-end machine learning pipeline, including data acquisition, model training, hyper-parameter tuning, post-processing of predictions, and model testing. ◦ Evaluated the model performance on Amazon Prime data and achieved 85% accuracy.	06-09/2019
Research Intern, Symantec Research Labs (NortonLifeLock) ◦ Proposed a federated machine learning framework that coordinates the training process of local nodes to jointly learn some desired model, while respecting the local data privacy during the communication between nodes. ◦ Published a paper in the International Joint Conference on Neural Networks (IJCNN).	05-08/2018

Publication

Superscript [★] for alphabetic author order.

Yuzhe Ma, Jon A. Sharp, Ruizhe Wang, Earlence Fernandes, Xiaojin Zhu. Adversarial Attacks on Kalman Filter-based Forward Collision Warning Systems. *In Submission* to **AAAI**, 2021.

Xuezhou Zhang, Shubham Bharti, **Yuzhe Ma**, Adish Singla, Xiaojin Zhu. The Teaching Dimension of Q-learning. Preprint in arXiv:2006.09324. *In Submission* to **AAAI**, 2021.

Xuezhou Zhang, **Yuzhe Ma**, Adish Singla. Task-agnostic Exploration in Reinforcement Learning. In The 34th Conference on Neural Information Processing Systems (**NeurIPS**), 2020.

Xuezhou Zhang, **Yuzhe Ma**, Adish Singla, Xiaojin Zhu. Adaptive Reward-Poisoning Attacks against Reinforcement Learning. In The 37th International Conference on Machine Learning (**ICML**), 2020.

Yuzhe Ma, Xuezhou Zhang, Wen Sun, Xiaojin Zhu. Policy Poisoning in Batch Reinforcement Learning and Control. In The 33rd Conference on Neural Information Processing Systems (**NeurIPS**), 2019.

Yuzhe Ma, Xiaojin Zhu, Justin Hsu. Data Poisoning against Differentially-Private Learners: Attacks and Defenses. In The 28th International Joint Conference on Artificial Intelligence (**IJCAI**), 2019.

Yufei Han, **Yuzhe Ma**, Chris Gates, Kevin Roundy, Yun Shen. Collaborative and Privacy-Preserving Machine Teaching via Consensus Optimization. In International Joint Conference on Neural Networks (**IJCNN**), 2019.

Kwang-Sung Jun[★], Lihong Li[★], **Yuzhe Ma**[★], Xiaojin Zhu[★]. Adversarial Attacks on Stochastic Bandits. In The 32nd Conference on Neural Information Processing Systems (**NeurIPS**), 2018.

Yuzhe Ma, Kwang-Sung Jun, Lihong Li, Xiaojin Zhu. Data Poisoning Attacks in Contextual Bandits. In The 9th Conference on Decision and Game Theory for Security (**GameSec**), 2018.

Ayon Sen, Scott Alfeld, Xuezhou Zhang, Ara Vartanian, **Yuzhe Ma**, Xiaojin Zhu. Training Set Camouflage. In The 9th Conference on Decision and Game Theory for Security (**GameSec**), 2018.

Yuzhe Ma, Robert Nowak, Philippe Rigollet, Xuezhou Zhang, Xiaojin Zhu. Teacher Improves Learning by Selecting a Training Subset. In The 21st International Conference on Artificial Intelligence and Statistics (**AISTATS**), 2018.

Yuzhe Ma, Kun He, John Hopcroft, Pan Shi. Neighbourhood-Preserving Dimension Reduction via Localised Multidimensional Scaling. In Theoretical Computer Science (**TCS**), 2017.

Yuzhe Ma, Kun He, John Hopcroft, Pan Shi. Nonlinear Dimension Reduction by Local Multidimensional Scaling. In The 10th International Frontiers of Algorithmics Workshop (**FAW**), 2016.

Yuzhe Ma, Kun He, Leihua Qin, Yan Wang. A Primary Research on Overlapping Community Detection. In The 32nd National Conference on Theoretical Computer Science (**NCTCS**), 2014.

Yun-Shiuan Chuang, Xuezhou Zhang, **Yuzhe Ma**, Mark K. Ho, Joseph L. Austerweil, Xiaojin Zhu. Using Machine Teaching to Investigate Human Assumptions when Teaching Reinforcement Learners. Preprint in arXiv:2009.02476.

Honors and Awards

Student Travel Award, NeurIPS	2019
Top 50% Reviewer, NeurIPS	2019
Student Travel Award, GameSec	2018
Honorarium Award, GameSec Special Track	2018
Student Travel Award, AISTATS	2018
UW CS Summer Research Award	2017
Outstanding Bachelor Thesis Award	2016
CCF Outstanding Undergraduate Award	2015
Outstanding Student Leader Award	2014
China National Scholarship Award	2013

Academic Service

Program Committee: AAAI21, ACML20, AAAI20, ACML19, AAAI19

Conference Reviewer: AISTATS21, ICLR21, NeurIPS20, ICML20, AISTATS20, NeurIPS19, ICML19, AISTATS19

Journal Reviewer: TON, TPAMI, IEEE Access, Machine Learning

Student Volunteer: AISTATS18

Skills & Expertise

Programming Skills: Python, Pytorch, SQL, Matlab, C, C++, R, AMPL, Verilog

Machine Learning: Multi-Armed Bandit, Reinforcement Learning, Deep Learning, Differential Privacy